



Transport Layer: Introduction to Sockets

Re-look at the stack

- Headers are “peeled” as you go up the stack
- Headers are added as you go down the stack.



Layering, Revisited




- Lowest level end-to-end protocol
 - Transport header only read by source and destination
 - Routers view transport header as payload
 - Each packet has a Maximum Segment Size (MSS)

Transport layer: TCP

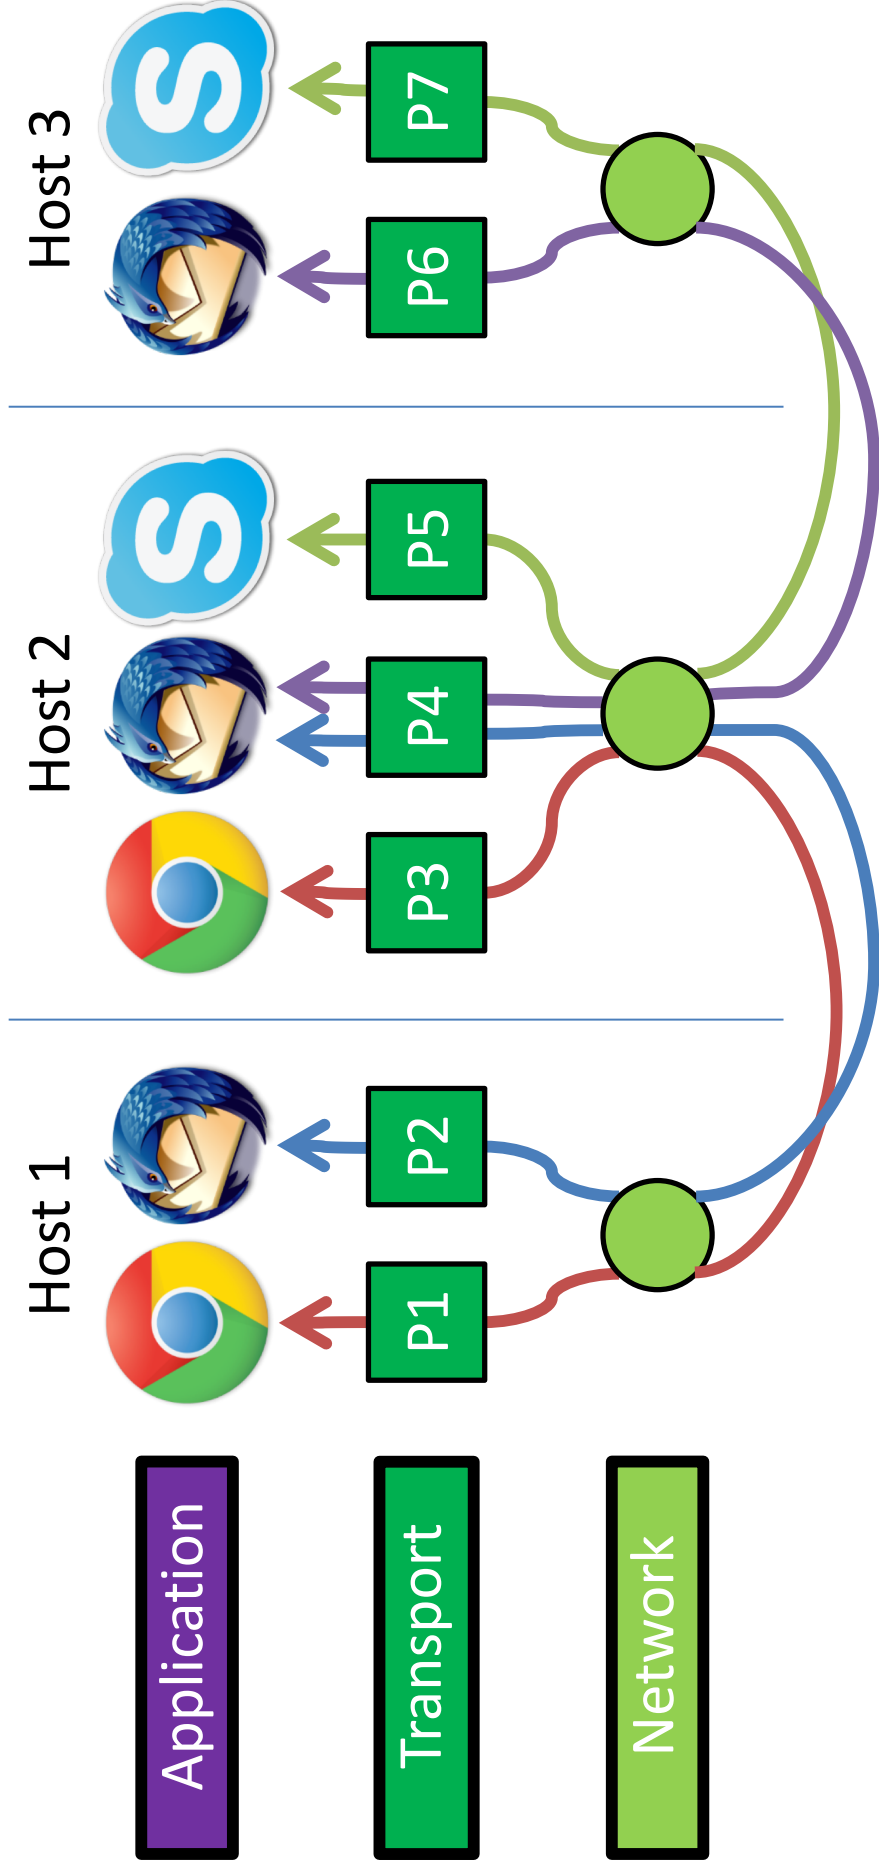
- Transport layer roles
 - “End-to-End” abstraction
 - De-multiplexing

What is de-multiplexing?

- Clients run many applications at the same time
 - Who to deliver packets to?
- Insert Transport Layer to handle demultiplexing using ports
- Rather than IP address, the end point has an IP address and a port.



Demultiplexing Traffic



Endpoints identified by `<src_ip, src_port, dest_ip, dest_port>`

Two types of Transport Protocol

- Transmission Control Protocol (TCP)
 - Connection oriented
 - Provides an “in-order delivery” abstraction
 - Masks unreliability.
- User Datagram Protocol (UDP)
 - Connection less
 - No guarantees of in-order delivery.
 - Does not mask unreliability.

A deep dive into Sockets

Socket programming with TCP

- client must contact server**
- server process must first be running
 - server must have created socket (door) that welcomes client's contact

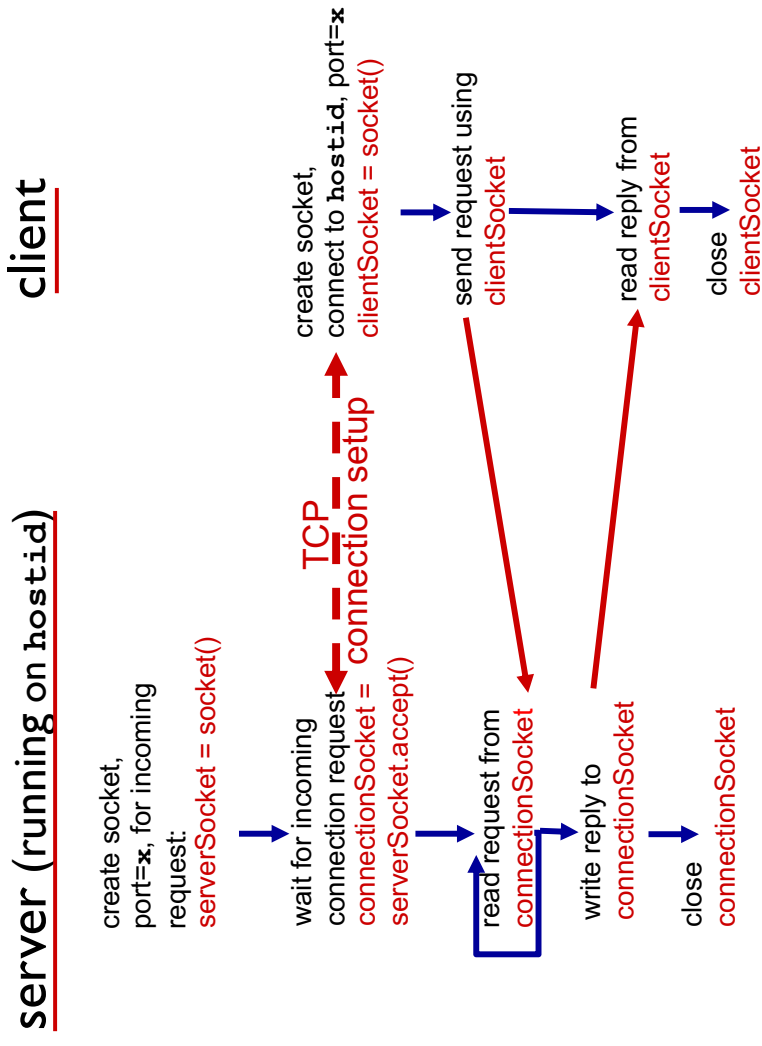
client contacts server by:

- Creating TCP socket, specifying IP address, port number of server process
- **when client creates socket:** client TCP establishes connection to server TCP

- when contacted by client, **server TCP creates new socket** for server process to communicate with that particular client
 - allows server to talk with multiple clients
 - source port numbers used to distinguish clients (more in Chap 3)
- application viewpoint:**

TCP provides reliable, in-order byte-stream transfer (“pipe”) between client and server

Client/server socket interaction: TCP



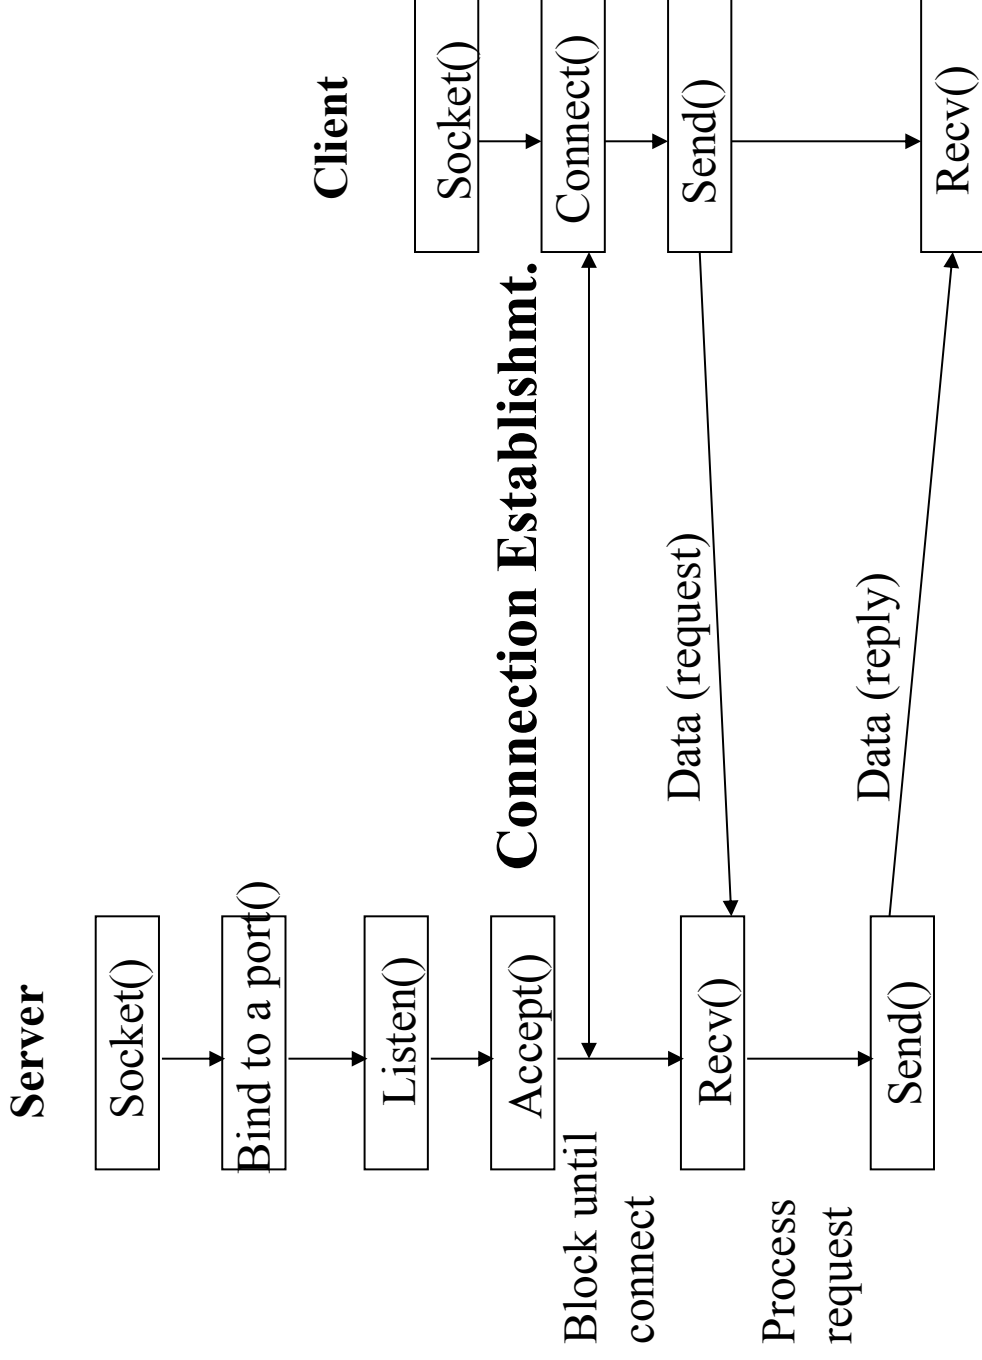
To make a TCP server

- Create a socket with the *socket()* system call.
- Bind the socket to an address using the *bind()* system call. For a server socket on the Internet, an address consists of a port number on the host machine.
- Listen for connections with the *listen()* system call.
- Accept a connection with the *accept()* system call. This call typically blocks until a client connects with the server.
- Send and receive data using the *read()* and *write()* system calls.

To make a TCP client

- Creates a socket with a `socket()` call
- Bind the socket to a newly created port
- Assume you already know the IP address and port number of server
- Connect to the server using a `connect()` call
- Start sending and receiving data.

TCP connection



Handling multiple connection

- To allow the server to handle multiple simultaneous connections, we make the following changes in the above code
- Put the *accept* statement and the following code in an infinite loop.
- After a connection is established, call *fork()* to create a new process.
- Do your processing in the new thread, the accept statement can continue to accept messages

Example app: TCP client

Python TCPClient

```
from socket import *
serverName = 'servername'
serverPort = 12000
clientSocket = socket(AF_INET, SOCK_STREAM)
clientSocket.connect((serverName, serverPort))
sentence = raw_input('Input lowercase sentence:')
clientSocket.send(sentence.encode())
modifiedSentence = clientSocket.recv(1024)
print ('From Server:', modifiedSentence.decode())
clientSocket.close()
```

create TCP socket for server, remote port 12000

No need to attach server name, port

Example app: TCP server

Python TCP Server

```
from socket import *
serverPort = 12000
serverSocket = socket(AF_INET, SOCK_STREAM)
serverSocket.bind(('', serverPort))
serverSocket.listen(1)
print 'The server is ready to receive'
while True:
    connectionSocket, addr = serverSocket.accept()

    sentence = connectionSocket.recv(1024).decode()
    capitalizedSentence = sentence.upper()
    connectionSocket.send(capitalizedSentence.
                           encode())
    connectionSocket.close()
```

create TCP welcoming socket →

server begins listening for incoming TCP requests →

loop forever →

server waits on accept() for incoming requests, new socket created on return →

read bytes from socket (but not address as in UDP) →

close connection to this client (but *not* welcoming socket) →