

Routing

arunab, SBU // CSE 310, Spring
2019: Intra-Domain Routing

Network layer: data plane, control plane

Data plane (Forwarding)

- local, per-router function
- determines how datagram arriving on router input port is forwarded to router output port
- forwarding function

Control plane (Routing)


- network-wide logic
- determines how datagram is routed among routers along end-end path from source host to destination host
- two control-plane approaches:
 - *traditional routing algorithms*: implemented in routers
 - *software-defined networking (SDN)*: implemented in (remote) servers

Routing protocols

Routing protocol goal: determine “good” paths (equivalently, routes), from sending hosts to receiving host, through network of routers

- path: sequence of routers packets will traverse in going from given initial source host to given final destination host
- “good”: least “cost”, “fastest”, “least congested”

Graph abstraction of the network




graph: $G = (N, E)$

N = set of routers = $\{ u, v, w, x, y, z \}$

E = set of links = $\{ (u,v), (u,x), (v,x), (v,w), (x,w), (x,y), (w,y), (w,z), (y,z) \}$

aside: graph abstraction is useful in other network contexts, e.g., P2P, where N is set of peers and E is set of TCP connections

Graph abstraction: costs



$c(x, x') = \text{cost of link } (x, x')$
e.g., $c(w, z) = 5$

cost could always be l, or
inversely related to bandwidth,
or inversely related to
congestion

cost of path $(x_1, x_2, x_3, \dots, x_p) = c(x_1, x_2) + c(x_2, x_3) + \dots + c(x_{p-1}, x_p)$

key question: what is the least-cost path between u and z ?
routing algorithm: algorithm that finds that least cost path

Routing algorithm classification

Q: global or decentralized information?

global:

- all routers have complete topology, link cost info

- “link state” algorithms

decentralized:

- router knows physically-connected neighbors, link costs to neighbors
- iterative process of computation, exchange of info with neighbors

- “distance vector” algorithms

Q: static or dynamic?

static:

- routes change slowly over time

dynamic:

- routes change more quickly
 - periodic update
 - in response to link cost changes

Analytical Background on Route Optimization

- **Network model:**
 - Network is a graph of nodes (routers) and links.
 - Each link k has a capacity: $c(k)$.
 - Traffic matrix: $M(i,j)$ = rate of traffic from node i to node j .
 - Proportion of traffic from node i to node j traversing link k : $P(i,j,k)$. If only one path from i to j , then $P(i,j,k)$ is either 1 or 0, depending on whether link k is on the path.
- **Classical algorithmic problem.**

In practice: Distributed, Dynamic Routing Protocols

- **Distributed** because in a dynamic network, no single, centralized node “knows” the whole “state” of the network.
- **Dynamic** because routing must respond to “state” changes in the network for efficiency.
- **Two types of protocols: Link State and Distance Vector.**
 - Link State uses Dijkstra’s shortest path algorithm, but makes up a distributed version.
 - Distance Vector uses a distributed version of Bellman Ford algorithm.

Routing uses standard shortest path algorithms

- **Derived from classic algorithms (e.g., Prims)**
 - Dijkstra's shortest path
 - Bellman-Ford
- **The key idea is to relax the distance as more information is provided**

Link State versus Distance Vector

- **Link state**
 - Send information to everyone in the network
 - Each will compute the shortest path
- **Distance vector**
 - Send information to neighbors only
 - Each neighbor determines shortest path based on one-hop information

Link State Protocol

- Each node “floods” the network with **link state packets (LSP)** describing the cost of its own (outgoing) links.
- Each node maintains a **LSP database** of all LSPs it received.
 - Only the recent most LSP is maintained for a link.
 - The LSP database describes this node’s view of the “state” of the network.
 - It is expected that all nodes “see” the same state (but of course not guaranteed)

Flooding Mechanism

- Flooding is a basic routing service. Used by many protocols in some form.
- The originator generates LSPs periodically, or when some link costs changes significantly.
 - The originator transmits LSP on all its interfaces.

A link-state routing algorithm

Dijkstra's algorithm


- net topology, link costs known to all nodes
 - accomplished via “link state broadcast”
- computes least cost paths from one node (‘source’) to all other nodes
 - gives *forwarding table* for that node

notation:


- $c(x,y)$: link cost from node x to y ; $= \infty$ if not direct neighbors
- $D(v)$: current value of cost of path from source to dest. v
- $p(v)$: predecessor node along path from source to v
- N' : set of nodes whose least cost path definitively known

Dijkstra's algorithm

```
1 Initialization:
2    $N' = \{u\}$ 
3   for all nodes  $v$ 
4     if  $v$  adjacent to  $u$ 
5       then  $D(v) = c(u,v)$ 
6     else  $D(v) = \infty$ 
7
8 Loop
9   find  $w$  not in  $N'$  such that  $D(w)$  is a minimum
10  add  $w$  to  $N'$ 
11  update  $D(v)$  for all  $v$  adjacent to  $w$  and not in  $N'$  :
12     $D(v) = \min(D(v), D(w) + c(w,v))$ 
13    /* new cost to  $v$  is either old cost to  $v$  or known
14       shortest path cost to  $w$  plus cost from  $w$  to  $v$  */
15 until all nodes in  $N'$ 
```



Link State (using Dijkstra's shortest path)



LSP database on
a node

A - B 1

A - D 3

B - E 4

.....

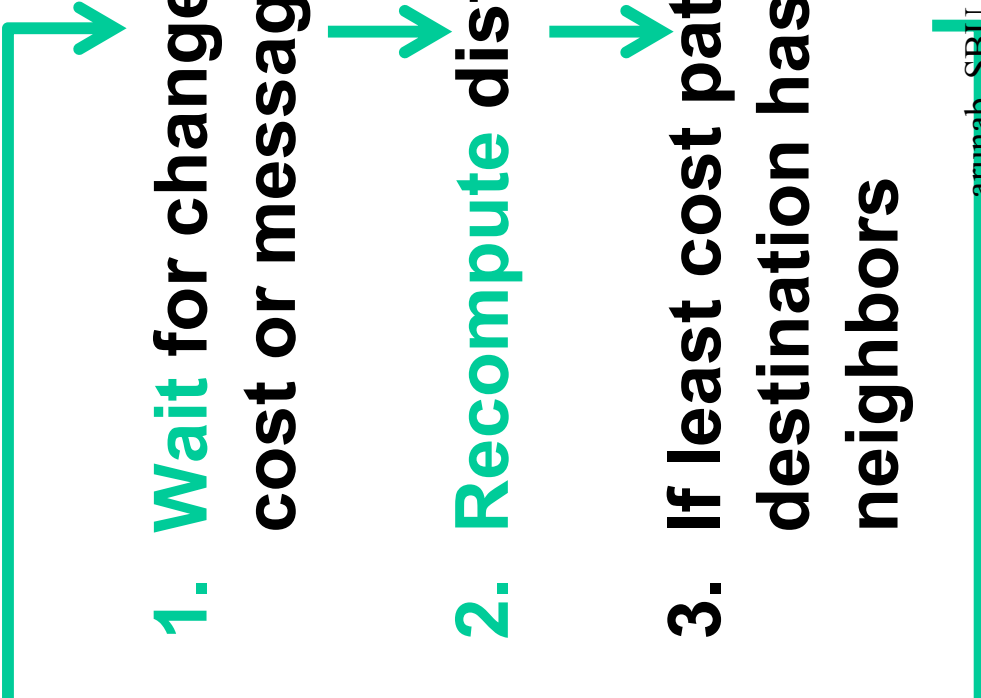
.....

- Important property of link state:
Every node has “all” information
about every other node.
- In class example.


Distance Vector Routing

- **What is a distance vector?**
 - Current best known cost to reach a destination
- **Idea: exchange vectors among neighbors to learn about lowest cost paths**

Distance Vector Routing Algorithm

1. **Wait** for change in local link cost or message from neighbor
 2. **Recompute** distance table
 3. If least cost path to any destination has changed, **notify** neighbors
- 

Distance Vector Initialization



Node A

Dest.	Cost	Next
B	2	B
C	7	C
D	∞	

Node B

Dest.	Cost	Next
A	2	A
C	1	C
D	3	D

Node C

Dest.	Cost	Next
A	7	A
B	1	B
D	1	D

Node D


Dest.	Cost	Next
A	∞	
B	3	B
C	1	C

1. Initialization:
2. for all neighbors V
 - do
3. if V adjacent to A
 - 4. $D(A, V) = c(A, V)$;
5. else
6. $D(A, V) = \infty$;
- ...

Distance Vector (DV) Routing

- See more details in class notes (loaded to resources section)
- Unlike link state, DV only knows the shortest path route to a destination from its neighbors.
- Distributed variation of Bellman-Ford algorithm.
- # of rounds to converge is roughly the length of the network.

Good news travels fast



Link Cost Changes,
Algorithm Starts

Node B

D	C	N
A	4	A
C	1	B

D	C	N
A	1	A
C	1	B

D	C	N
A	1	A
C	1	B

Node C

D	C	N
A	5	B
B	1	B



D	C	N
A	5	B
B	1	B

D	C	N
A	2	B
B	1	B

Source: SDBU // CS2210, Spring

2019: Intra-Domain Traffic

Bad news travels slowly: Count to Infinity Problem



Count to infinity problem

- Because of cycles
- There are some ways to improve this, we will not cover them

Distance Vector (RIP) vs. Link State (OSPF)

- RIP uses UDP to exchange information
- OSPF uses IP packets to exchange information

Which is best?
In practice, it depends.
In general, link state is more popular.